

```
In [1]: import pandas as pd

import warnings
warnings.filterwarnings('ignore')

# Load dataset
df = pd.read_csv("D:\\mcdonalds.csv")
```

```
In [2]: df.head()
```

```
Out[2]:
```

	yummy	convenient	spicy	fattening	greasy	fast	cheap	tasty	expensive	healthy	disgusting	Like	Age	VisitFrequency	Gender
0	No	Yes	No	Yes	No	Yes	Yes	No	Yes	No	No	-3	61	Every three months	Female
1	Yes	Yes	No	Yes	Yes	Yes	Yes	Yes	Yes	No	No	+2	51	Every three months	Female
2	No	Yes	Yes	Yes	Yes	Yes	No	Yes	Yes	Yes	No	+1	62	Every three months	Female
3	Yes	Yes	No	Yes	Yes	Yes	Yes	Yes	No	No	Yes	+4	69	Once a week	Female
4	No	Yes	No	Yes	Yes	Yes	Yes	No	No	Yes	No	+2	49	Once a month	Male

```
In [3]: # Drop any rows with missing values
df.dropna(inplace=True)

# Ensure the 'Age' column is numeric
df['Age'] = pd.to_numeric(df['Age'], errors='coerce')


# Drop rows with invalid Age
df.dropna(subset=['Age'], inplace=True)
```

```
In [4]: # Convert Yes/No to 1/0 for all binary columns (excluding 'Like')
binary_columns = ['yummy', 'convenient', 'spicy', 'fattening', 'greasy', 'fast', 'cheap', 'tasty', 'expensive', 'healthy', 'disgusting']
df[binary_columns] = df[binary_columns].apply(lambda x: x.map({'Yes': 1, 'No': 0}))
```

```
In [5]: from sklearn.cluster import KMeans
from statsmodels.graphics.mosaicplot import mosaic
import matplotlib.pyplot as plt

# Apply KMeans clustering to the binary columns (excluding 'Like')
kmeans = KMeans(n_clusters=4, random_state=42, n_init=10)
df['Segment'] = kmeans.fit_predict(df[binary_columns])


# Mosaic plot for Segment vs Gender
mosaic(df, ['Segment', 'Gender'])
plt.title('Segment vs Gender')
plt.show()
```




```
In [6]: import seaborn as sns

# Boxplot of Age by Segment
sns.boxplot(x='Segment', y='Age', data=df)
plt.title('Age by Segment')
plt.show()

# Histograms of Age by Segment
g = sns.FacetGrid(df, col="Segment", col_wrap=2)
g.map(plt.hist, "Age", bins=15)
plt.show()
```





```
In [7]: from scipy.stats import chi2_contingency

# Cross-tabulation and chi-square test for Gender vs Segment
gender_table = pd.crosstab(df['Segment'], df['Gender'])
chi2, p, dof, expected = chi2_contingency(gender_table)
```

```
print("Chi-square p-value:", p)
```

Chi-square p-value: 8.697992697736528e-07

```
In [8]: import statsmodels.api as sm
        from statsmodels.formula.api import ols
        from statsmodels.stats.multicomp import pairwise_tukeyhsd

        # ANOVA to compare Age across Segments
        model = ols('Age ~ C(Segment)', data=df).fit()
        anova_table = sm.stats.anova_lm(model, typ=2)
        print(anova_table)

        # Tukey's HSD for pairwise comparison of Age between Segments
        tukey = pairwise_tukeyhsd(endog=df['Age'], groups=df['Segment'], alpha=0.05)
        print(tukey)
```

	sum_sq	df	F	PR(>F)
C(Segment)	16966.145572	3.0	29.616809	1.390447e-18
Residual	276689.098750	1449.0	NaN	NaN

Multiple Comparison of Means - Tukey HSD, FWER=0.05

```
=====
```

group1	group2	meandiff	p-adj	lower	upper	reject
0	1	-4.6085	0.0001	-7.3363	-1.8806	True
0	2	0.937	0.8599	-2.121	3.995	False
0	3	-8.2242	0.0	-11.2511	-5.1973	True
1	2	5.5455	0.0	3.0423	8.0487	True
1	3	-3.6158	0.001	-6.0808	-1.1507	True
2	3	-9.1612	0.0	-11.9873	-6.3352	True

```
-----
```

```
In [9]: from sklearn.linear_model import LogisticRegression
        from sklearn.model_selection import train_test_split

        # Define target variable: Predict 'Like'
        X = pd.get_dummies(df[['Age', 'Gender']], drop_first=True) # Feature matrix (Age, Gender)
        y = df['Like'] # Target variable (Like)

        # Split the data into training and testing sets
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)

# Train logistic regression model
log_reg = LogisticRegression(max_iter = 1000)
log_reg.fit(X_train, y_train)

# Print accuracy
print("Logistic regression accuracy:", log_reg.score(X_test, y_test))
```

Logistic regression accuracy: 0.15825688073394495

```
In [10]: from sklearn.linear_model import LogisticRegression

# Define features and target
X = pd.get_dummies(df[['Age', 'Gender']], drop_first=True)
y = df['Segment']

# Train multinomial logistic regression model
multi_model = LogisticRegression(multi_class='multinomial', solver='lbfgs', max_iter=1000)
multi_model.fit(X, y)

# Print accuracy
print("Multinomial Logistic Regression Accuracy:", multi_model.score(X, y))
```

Multinomial Logistic Regression Accuracy: 0.39022711631108054

```
In [11]: from sklearn.tree import DecisionTreeClassifier, plot_tree

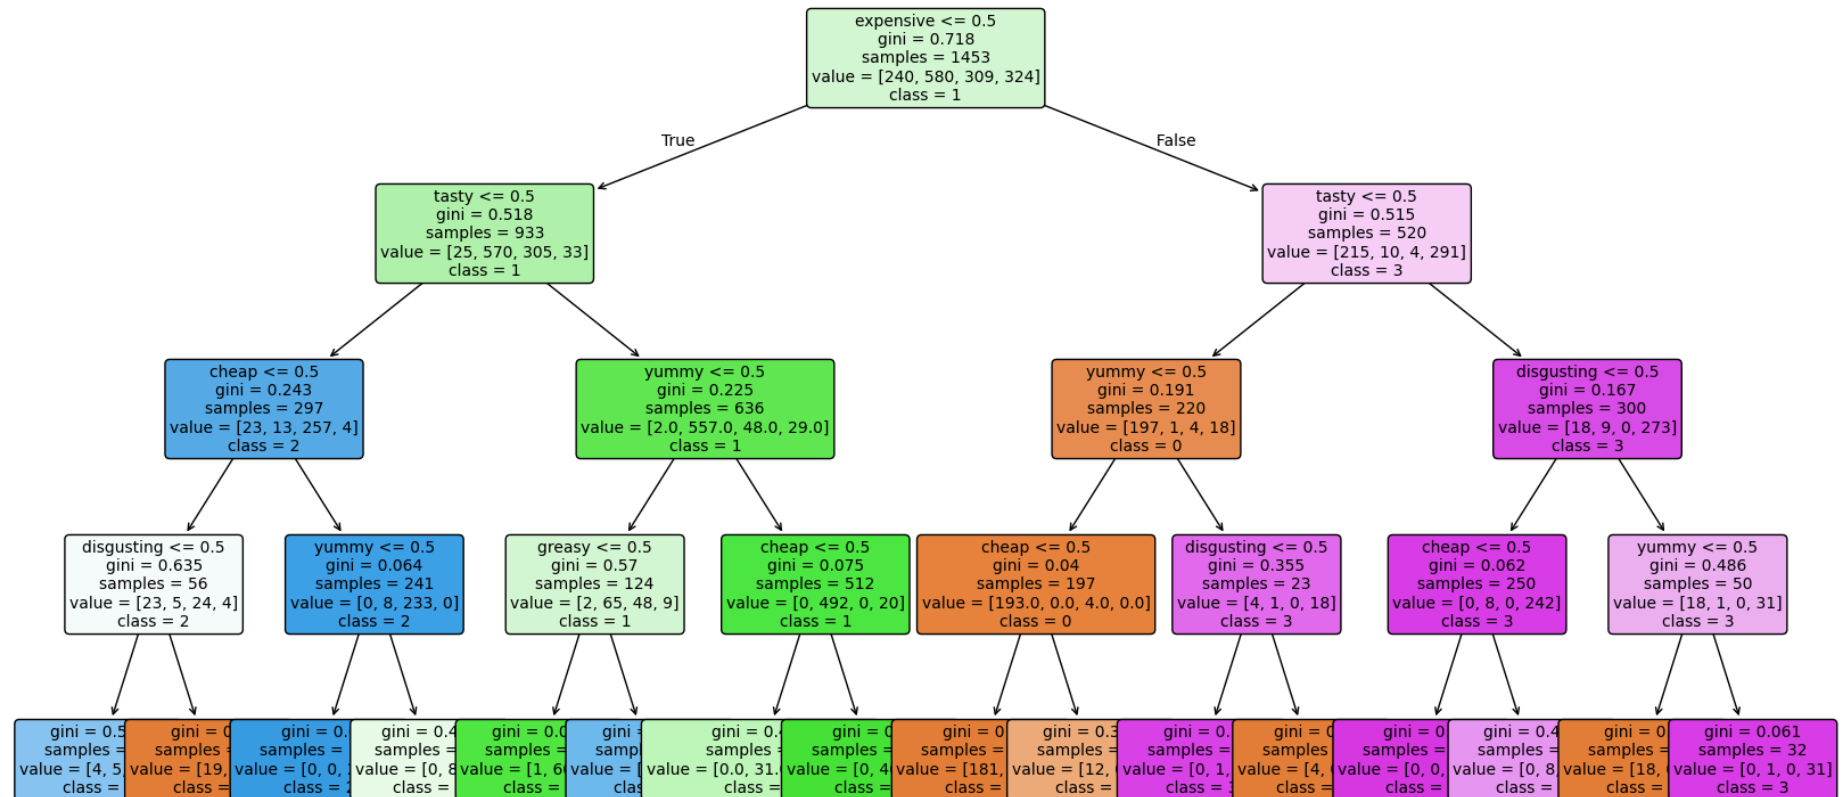
X = pd.get_dummies(df.drop(columns=['Segment', 'Like']), drop_first=True)
y = df['Segment']

# Train the model
tree_model = DecisionTreeClassifier(max_depth=4, random_state=42)
tree_model.fit(X, y)

# Visualize the decision tree
plt.figure(figsize=(20, 10))
plot_tree(tree_model,
          filled=True,
          feature_names=X.columns,
          class_names=[str(i) for i in sorted(df['Segment'].unique())],
```

```
rounded=True,
fontsize=10)
plt.title("Decision Tree for Customer Segmentation")
plt.show()
```

Decision Tree for Customer Segmentation



In [ ]: